

Kafka Consumer Properties



Kafka Consumer Configs

- 1) bootstrap.server
- 2) client.id
- 3) key.deserializer
- 4) value.deserializer
- 5) group.id
- 6) fetch.min.bytes

Kindly explain in Short with respect to above. What are the different types of client in Kafka



The different types of clients in Kafka include:

1. **Producer:** Sends messages to Kafka topics. It uses configurations like `bootstrap.servers` to connect to the cluster and sends data to the specified topic partitions.
2. **Consumer:** Reads messages from Kafka topics. It relies on configurations like those in the image:
 - `bootstrap.server`: Kafka broker to connect to.
 - `group.id`: Consumer group identifier for load balancing and fault tolerance.
 - `key.deserializer` & `value.deserializer`: Convert Kafka messages (key/value) from bytes to usable data types.
 - `fetch.min.bytes`: Controls the minimum data size the broker sends to the consumer.
3. **Stream Processor:** Processes and transforms data streams from input topics and writes results to output topics.
4. **Admin Client:** Manages Kafka resources like topics, configurations, and partitions programmatically.

These clients enable Kafka's flexible and scalable use in data pipelines and real-time applications.



BOOTSTRAP.SERVER

bootstrap.servers

A list of host/port pairs to use for establishing the initial connection to the Kafka cluster. The client will make use of all servers irrespective of which servers are specified here for bootstrapping—this list only impacts the initial hosts used to discover the full set of servers. This list should be in the form host1:port1,host2:port2,.... Since these servers are just used for the initial connection to discover the full cluster membership (which may change dynamically), this list need not contain the full set of servers (you may want more than one, though, in case a server is down).

TYPE: list
DEFAULT: ""
VALID VALUES: non-null string
IMPORTANCE: high

CLIENT.ID

client.id

An id string to pass to the server when making requests. The purpose of this is to be able to track the source of requests beyond just ip/port by allowing a logical application name to be included in server-side request logging.

TYPE: string
DEFAULT: ""
VALID VALUES:
IMPORTANCE: low

Kindly explain above with respect to bootstrap.properties and Client.id with respect to Consumer in Kafka



Explanation of `bootstrap.servers` and `client.id` in Kafka Consumer Configuration:

1. `bootstrap.servers`

- **Purpose:** Specifies a list of host:port pairs for the Kafka cluster that the client (consumer) should connect to initially.
- **Use in `bootstrap.properties`:**

- This property is required in the consumer configuration file.
- It helps the consumer establish an initial connection to the Kafka cluster to fetch metadata about the cluster and available topics.
- Format: ``host1:port1,host2:port2,...``
- Example:

```
properties
```

```
bootstrap.servers=localhost:9092,broker2:9093
```

- **Note:** Only the initial connection uses this list; Kafka dynamically discovers the rest of the brokers.

2. ``client.id``

- **Purpose:** Provides an identifier for the consumer client. It allows logical grouping or tracking of requests made to Kafka servers. This is helpful for server-side logging and debugging.
- **Use in ``bootstrap.properties``:**
 - Adding this property makes it easier to trace which client made specific requests in Kafka server logs.
 - Example:

```
properties
```

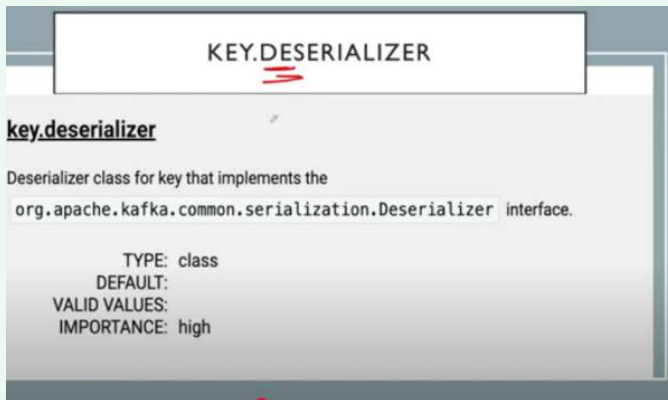
```
client.id=my-consumer-client
```

- Default: If not specified, Kafka assigns a default identifier.
- **Importance:** Low, but useful for debugging or monitoring.

Summary:

In the consumer context:

- ``bootstrap.servers`` is crucial for connecting to Kafka, ensuring the consumer knows where to start.
- ``client.id`` is optional but aids in identifying consumers during operations and monitoring.



Kindly explain above and What is the default value deSerializer in Consumer



Explanation of `key.deserializer`

What is `key.deserializer`?

- It specifies the class used by the Kafka consumer to **deserialize the keys** of messages received from Kafka topics.
- Deserialization is the process of converting the byte array (in which Kafka stores keys) into a usable Java object.

Details:

- The `key.deserializer` must implement the `org.apache.kafka.common.serialization.Deserializer` interface.
- This deserializer converts the key from its serialized (byte array) form back into the object form expected by the application.

Configuration in Consumer:

- This property is mandatory for consuming messages.
- Example in `bootstrap.properties`:

```
properties

key.deserializer=org.apache.kafka.common.serialization.StringDeserializer
```

Default Value for Deserializer

There is **no default value** for `key.deserializer`. It must be explicitly set by the user. However, the most commonly used deserializers are:

- **StringDeserializer:** Converts the byte array into a String. Often used when keys are simple text.

```
properties

key.deserializer=org.apache.kafka.common.serialization.StringDeserializer
```

- **ByteArrayDeserializer:** If the key is stored as raw bytes, this deserializer is used.

```
properties

key.deserializer=org.apache.kafka.common.serialization.ByteArrayDeserializer
```

Importance:

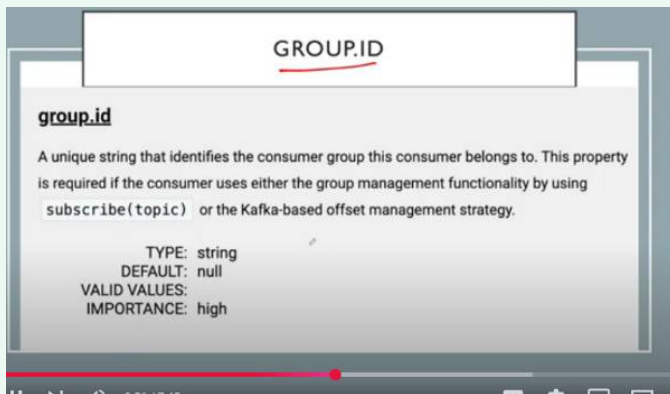
- If this property is not configured correctly, the consumer cannot understand the keys of the messages, resulting in deserialization errors.
- The choice of deserializer must match the serializer used by the Kafka producer for the same topic.

Example Usage:

If the producer serialized the key as a String, the consumer should use `StringDeserializer`:

```
java
```

```
key.deserializer=org.apache.kafka.common.serialization.StringDeserializer
```



Explain the groupid wrt to consumer in Kafka



Explanation of `group.id` in Kafka Consumer

What is `group.id`?

- `group.id` is a unique string that identifies the **consumer group** a particular consumer belongs to.
- **Consumer Group**: A collection of consumers that coordinate to read data from a Kafka topic.

Key Points:

1. Partition Assignment:

- Within a consumer group, Kafka ensures that each partition of a topic is consumed by **only one consumer** in the group.
- This allows parallel processing of messages across multiple consumers.

2. Group Coordination:

- Kafka uses the ``group.id`` to keep track of offsets (the position of the last consumed message) for each consumer in the group.
- This ensures that consumers can pick up where they left off in case of a restart.

3. Shared Workload:

- Multiple consumers in the same group can share the workload by consuming different partitions of a topic.
- If multiple consumers share the same ``group.id``, the partitions are dynamically assigned to them.

4. Independent Consumption:

- Consumers with **different** ``group.id``s consume data independently of each other, even if they are reading from the same topic.

Usage in Consumer Configuration:

- Example in ``bootstrap.properties``:

```
properties

group.id=my-consumer-group
```

- This ensures that all consumers sharing the same ``group.id`` are treated as part of the same group.

Importance of ``group.id``:

- Without a ``group.id``, consumers cannot use Kafka's offset management functionality.
- It is mandatory when using the ``subscribe()`` method to consume messages.
- **Default:** Null (must be explicitly set).

Example Scenarios:

1. Load Balancing:

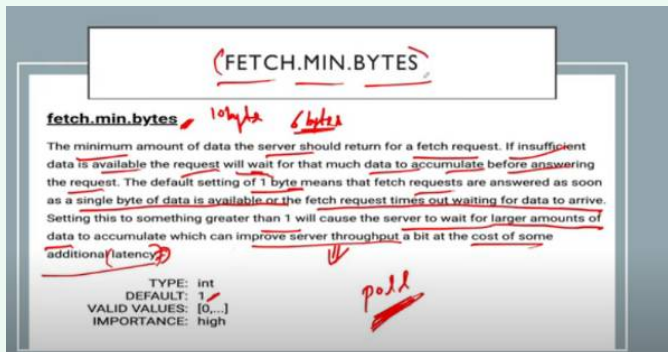
- Topic ``my-topic`` has 4 partitions.
- Consumer group with ``group.id=my-group`` has 4 consumers, and each consumer will consume from one partition.

2. Independent Consumers:

- Two groups with ``group.id=my-group`` and ``group.id=other-group`` both consume from ``my-topic``.
- Each group will get all messages independently.

Summary

The ``group.id`` is critical for enabling Kafka's consumer group functionality, including offset management, workload sharing, and fault tolerance. Proper configuration ensures efficient and coordinated message consumption.



Explain the concept of Fetch.MinBytes in details as shown in image. Kindly explain with example



Explanation of ``fetch.min.bytes`` in Kafka Consumer

What is ``fetch.min.bytes``?

- **Purpose:** This configuration controls the **minimum amount of data** that the Kafka broker will return for a fetch request from the consumer.
- If the available data is less than the specified value, the broker will **wait** for more data to accumulate before sending a response to the consumer.

Default Value:

- The default value is **1 byte**, meaning the broker responds as soon as at least 1 byte of data is available for the consumer.

Detailed Behavior:

1. Small Value (e.g., 1 byte):

- The consumer receives data as soon as it becomes available, even if it's a very small amount.
- This results in **low latency** but can cause **more frequent fetch requests** and higher server overhead.

2. Larger Value (e.g., 10 KB):

- The broker waits until at least 10 KB of data is available before sending it to the consumer.
- This improves **throughput** by reducing the number of fetch requests but adds **latency**, as the broker waits for more data to accumulate.

How It Works in Fetch Requests:

1. Fetch Request from Consumer:

- The consumer sends a request to the broker to fetch messages.
- The broker checks how much data is available in the partition.

2. Broker Response:

- If the available data meets or exceeds `fetch.min.bytes`, the broker immediately sends the data back.
- If not, the broker waits until enough data accumulates (or until the request times out, based on `fetch.max.wait.ms`).

Example:

Scenario 1: Small `fetch.min.bytes`

- Configuration:

```
properties

fetch.min.bytes=1
```

- Behavior:

- A topic has a partition with only 2 KB of data.
- The consumer fetches this data immediately.
- Result: **Low latency**, but frequent fetch requests (inefficient for throughput).

Scenario 2: Larger `fetch.min.bytes`

- Configuration:

```
properties

fetch.min.bytes=10 KB
```

- Behavior:
 - A topic has a partition with only 6 KB of data.
 - The broker waits for more data to accumulate until it reaches 10 KB or the fetch request times out.
 - Result: **Higher throughput**, but increased latency since the broker waits for more data.

When to Use Larger Values:

- Use larger values of `fetch.min.bytes` when:
 - You prioritize **throughput** over **latency**.
 - Your application processes large batches of data rather than individual records.
 - You want to reduce the number of fetch requests to optimize server performance.

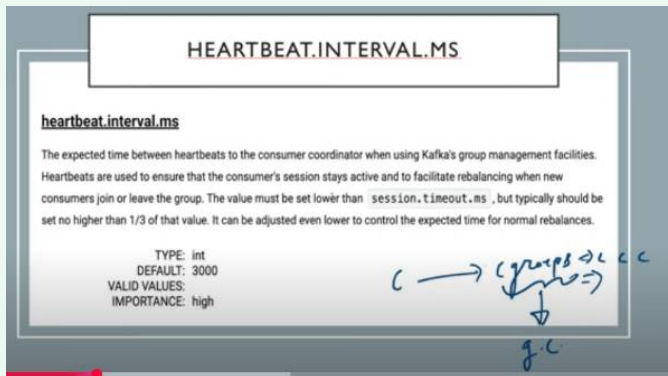
Key Considerations:

- **Tradeoff:**
 - A larger `fetch.min.bytes` value increases **throughput** but adds **latency**.
 - A smaller value reduces latency but increases fetch request overhead.
- **Timeout:**
 - If the broker cannot accumulate the specified amount of data before `fetch.max.wait.ms` expires, it will respond with whatever data is available.

Summary:

The `fetch.min.bytes`` property is a balancing tool between **latency** and **throughput**. The ideal value depends on your use case:

- For real-time applications: Use a smaller value for faster responses.
- For batch processing: Use a larger value for better efficiency.



Kindly explain the above image in details along with example



Explanation of `heartbeat.interval.ms`` in Kafka Consumer

What is `heartbeat.interval.ms``?

- It defines the expected time (in milliseconds) between successive heartbeat messages sent by a consumer to the **Kafka Group Coordinator**.
- Heartbeats are used to:
 1. Keep the consumer session **alive** within the consumer group.
 2. Facilitate **group rebalancing** when new consumers join or leave the group.

Detailed Behavior:

1. Consumer Group Management:

- When a consumer subscribes to a topic, it joins a **consumer group**.
- The Kafka Group Coordinator manages this group, ensuring active consumers are assigned partitions.

2. Heartbeats:

- Consumers send heartbeat messages periodically to inform the Group Coordinator that they are still active.
- If the consumer fails to send heartbeats within the configured ``session.timeout.ms``, it is considered dead, and its partitions are reassigned to other consumers in the group.

3. Relation to ``session.timeout.ms``:

- ``heartbeat.interval.ms`` should always be set **lower than** ``session.timeout.ms``.
- Typically, it is **1/3 of** ``session.timeout.ms`` to allow multiple attempts before the session times out.

Default Value:

- Default: ``3000 ms`` (3 seconds).

Example Configuration:

Scenario 1: Default Settings

- Configuration:

```
properties

heartbeat.interval.ms=3000
session.timeout.ms=10000
```

- Behavior:
 - Heartbeats are sent every 3 seconds.
 - If no heartbeat is received for 10 seconds, the consumer is considered dead, and its partitions are reassigned.

Scenario 2: Custom Settings

- Configuration:

```
properties
```

```
heartbeat.interval.ms=1000  
session.timeout.ms=3000
```

- Behavior:

- Heartbeats are sent every 1 second.
- If no heartbeat is received for 3 seconds, the consumer is marked as dead.
- This configuration ensures quicker detection of consumer failures but increases the network load due to frequent heartbeats.

Why is `heartbeat.interval.ms` Important?

1. Fault Tolerance:

- Ensures that inactive or failed consumers are detected promptly, enabling efficient partition reassignment.

2. Rebalancing:

- Facilitates smooth rebalancing of partitions when consumers join or leave the group.

3. Tradeoff:

- **Smaller interval:** Faster failure detection but higher network overhead.
- **Larger interval:** Reduced network load but slower detection of failed consumers.

Practical Example:

Real-Time Streaming Application:

- You have a consumer group processing a real-time data stream with 4 consumers.
- Configuration:

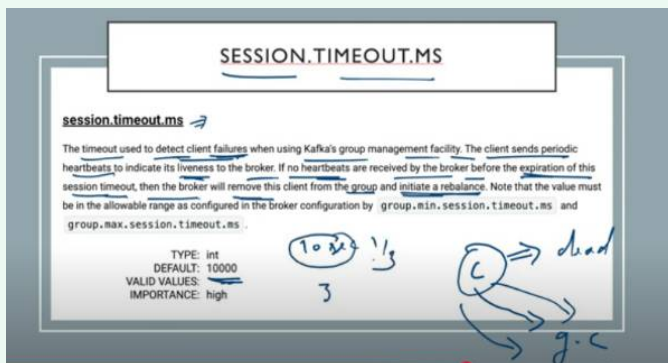
properties

```
heartbeat.interval.ms=2000
session.timeout.ms=6000
```

- Scenario:
 - Consumer A is active and sends heartbeats every 2 seconds.
 - If Consumer A crashes and fails to send heartbeats for 6 seconds, the Group Coordinator detects the failure and reassigns Consumer A's partitions to the remaining consumers.

Summary:

The `heartbeat.interval.ms` setting ensures the reliability and stability of Kafka consumer groups by maintaining active communication with the Group Coordinator. Adjusting it appropriately balances failure detection speed and network efficiency.



Kindly explain the above image in details



Explanation of `session.timeout.ms` in Kafka Consumer

What is `session.timeout.ms`?

- `session.timeout.ms` is the amount of time (in milliseconds) a Kafka broker waits to receive a **heartbeat** from a consumer before it considers the consumer **dead**.
 - If no heartbeat is received within this timeout:
 1. The broker **removes the consumer** from the group.
 2. It **initiates a rebalance** to reassign the consumer's partitions to other active consumers.
-

Key Details:

1. Heartbeats and Liveness:

- Consumers periodically send **heartbeats** to the broker to indicate they are alive.
- This property specifies the maximum time a consumer can go without sending a heartbeat before the broker takes action.

2. Group Rebalancing:

- When a consumer is removed due to a session timeout, Kafka performs a rebalance within the group, redistributing partitions among remaining consumers.

3. Broker Constraints:

- The value of `session.timeout.ms` must fall within the range specified by the broker:
 - `group.min.session.timeout.ms` (default: 6000 ms)
 - `group.max.session.timeout.ms` (default: 300000 ms)

4. Relation to `heartbeat.interval.ms`:

- `heartbeat.interval.ms` should always be set **lower than** `session.timeout.ms`, typically **1/3** of `session.timeout.ms` to allow multiple heartbeats to be sent before the session times out.
-

Default Value:

- Default: **10,000 ms** (10 seconds).
-

How It Works:

1. Normal Operation:

- The consumer sends heartbeats every `heartbeat.interval.ms` (e.g., 3 seconds).`
- If the broker receives heartbeats within the `session.timeout.ms` window, the consumer is considered alive.`

2. Failure Detection:

- If a consumer crashes or is unable to send heartbeats, the broker waits for `session.timeout.ms` before marking it as dead and reassigning its partitions.`

Example Configuration:

Scenario 1: Default Settings

- Configuration:

```
properties

session.timeout.ms=10000
heartbeat.interval.ms=3000
```

- Behavior:
 - Heartbeats are sent every 3 seconds.
 - If no heartbeat is received for 10 seconds, the broker considers the consumer dead and starts rebalancing.

Scenario 2: Custom Settings for Faster Detection

- Configuration:

```
properties

session.timeout.ms=6000
heartbeat.interval.ms=2000
```

- Behavior:
 - Heartbeats are sent every 2 seconds.
 - If no heartbeat is received for 6 seconds, the broker removes the consumer from the group.

Use Cases and Importance:

1. Detecting Dead Consumers:

- Ensures that failed or disconnected consumers are quickly removed from the group.
- Partitions handled by the dead consumer are reassigned to active consumers to maintain processing continuity.

2. Trade-Off:

- **Lower value:** Faster detection but increased risk of false positives (e.g., temporary network issues causing missed heartbeats).
- **Higher value:** Slower detection, but more tolerant of transient issues.

Practical Example:

Real-Time Application:

- You have a consumer group processing financial transactions.
- Configuration:

```
properties

session.timeout.ms=5000
heartbeat.interval.ms=1500
```

- Behavior:
 - Heartbeats are sent every 1.5 seconds.
 - If no heartbeat is received within 5 seconds (e.g., due to consumer crash), the broker reassigns the consumer's partitions to other group members.

Summary:

The `session.timeout.ms` setting is critical for detecting and handling consumer failures in a timely manner. It ensures that the Kafka broker maintains an active and balanced consumer group, facilitating high availability and fault tolerance. Adjust the value based on your application's failure detection and tolerance requirements.

MAX.PARTITION.FETCH.BYTES

max.partition.fetch.bytes

The maximum amount of data per-partition the server will return. Records are fetched in batches by the consumer. If the first record batch in the first non-empty partition of the fetch is larger than this limit, the batch will still be returned to ensure that the consumer can make progress. The maximum record batch size accepted by the broker is defined via `message.max.bytes` (broker config) or `max.message.bytes` (topic config). See `fetch.max.bytes` for limiting the consumer request size.

TYPE: int
DEFAULT: 1048576
VALID VALUES: [0,...]
IMPORTANCE: high

FETCH.MAX.BYTES

fetch.max.bytes

The maximum amount of data the server should return for a fetch request. Records are fetched in batches by the consumer, and if the first record batch in the first non-empty partition of the fetch is larger than this value, the record batch will still be returned to ensure that the consumer can make progress. As such, this is not an absolute maximum. The maximum record batch size accepted by the broker is defined via `message.max.bytes` (broker config) or `max.message.bytes` (topic config). Note that the consumer performs multiple fetches in parallel.

TYPE: int
DEFAULT: 52428800
VALID VALUES: [0,...]
IMPORTANCE: medium



Kindly explain the above image in details



Explanation of `max.partition.fetch.bytes` and `fetch.max.bytes` in Kafka Consumer

1. `max.partition.fetch.bytes`

- **Definition:**

- Specifies the maximum amount of data **per partition** that the Kafka broker will return in a single fetch request for a consumer.

- **Purpose:**

- Controls the size of data fetched **per partition** to optimize memory usage and ensure smoother data processing.
 - **Default Value:**
 - **1,048,576 bytes (1 MB).**
 - **Behavior:**
 - If a partition contains data larger than this limit, the broker ensures the first record batch in the partition is returned, even if it exceeds this size.
 - Ensures the consumer progresses even when large record batches exist.
 - **Example:**
 - If `max.partition.fetch.bytes=2 MB` and a topic has 4 partitions, the consumer can fetch up to 2 MB from each partition in a single request, potentially fetching up to 8 MB in total (4 partitions × 2 MB).
-

2. `fetch.max.bytes`

- **Definition:**
 - Specifies the maximum amount of data **for the entire fetch request** that the broker will return to the consumer.
 - **Purpose:**
 - Controls the overall size of data fetched in a single request to limit memory usage at the consumer level.
 - **Default Value:**
 - **52,428,800 bytes (50 MB).**
 - **Behavior:**
 - Even if `max.partition.fetch.bytes` allows more data, the total fetched data across all partitions is capped by `fetch.max.bytes`.
 - The broker might still return a batch that exceeds this size if it is the first batch in the partition and necessary for consumer progress.
 - **Example:**
 - If `fetch.max.bytes=10 MB` and the consumer requests data from 4 partitions, it can fetch up to 10 MB across all partitions combined, regardless of the `max.partition.fetch.bytes` value.
-

Relationship Between `max.partition.fetch.bytes` and `fetch.max.bytes`

1. Per Partition vs. Total Fetch:

- `max.partition.fetch.bytes` applies to each partition, while `fetch.max.bytes` applies to the sum of all partitions.

2. Interaction:

- If `max.partition.fetch.bytes` is set too high relative to `fetch.max.bytes`, only a limited number of partitions may contribute data to the fetch request.
- Example:
 - `max.partition.fetch.bytes=5 MB`
 - `fetch.max.bytes=10 MB`
 - Only two partitions can contribute 5 MB each in a single fetch request.

3. Priority:

- The consumer always fetches at least one record batch from each partition to ensure progress, even if these limits are exceeded temporarily.

Best Practices

1. Balance Memory and Performance:

- Adjust `max.partition.fetch.bytes` to avoid excessive memory usage for each partition.
- Configure `fetch.max.bytes` to control total memory usage and prevent overloading the consumer.

2. Use Case Scenarios:

- **High Throughput:** Set higher limits to fetch more data in each request, reducing fetch frequency.
- **Real-Time Applications:** Use lower limits to minimize latency and avoid large fetch delays.

Practical Configuration Example:

- **Scenario:** A consumer fetches data from a topic with 10 partitions.

```
properties
```

```
max.partition.fetch.bytes=2 MB  
fetch.max.bytes=10 MB
```

- **Behavior:**

- The consumer can fetch up to 2 MB from each partition, but the total data fetched will not exceed 10 MB.
- If three partitions have large record batches of 3 MB each:

- Only data from two partitions (up to 2 MB each) is fetched in one request.
-

Summary

- ``max.partition.fetch.bytes`` limits data fetched **per partition**, optimizing memory usage for large partitions.
- ``fetch.max.bytes`` limits total data fetched per request, ensuring the consumer doesn't fetch more than it can handle in memory.
- Proper configuration of these parameters is essential for balancing throughput, memory usage, and consumer performance.