

# Kafka Internal Producer



## OFFSETS



The records in the partitions are each assigned a sequential id number called the offset that uniquely identifies each record within the partition.

Three variations of offset

- 1) Log-end offset :- Offset of the last message written to a log/Partition.
- 2) Current offset :- Pointer to the last record that Kafka has already sent to a consumer in the most recent poll.
- 3) Committed offset :- Marking an offset as consumed is called committing an offset(Committed offset).

## INTERNALS WHEN PRODUCER SENDS MESSAGES



Kindly explain above with reference to image



The images depict concepts related to **Kafka offsets** and the **internals of message partitioning** when a producer sends messages to a Kafka topic.

### Image 1: Offsets in Kafka

Kafka uses offsets to identify messages uniquely within a partition. Here's a breakdown of the concepts:

### 1. Log-end Offset:

- This is the offset of the last message written to a partition.
- It's the boundary where new records will be appended.

### 2. Current Offset:

- Refers to the pointer to the last record that Kafka has already sent to a consumer during the most recent poll.
- The consumer uses this to know where to resume processing from.

### 3. Committed Offset:

- This is the offset that a consumer has explicitly marked as "processed" or "consumed".
- It's used to ensure that, in case of a failure, the consumer can restart from this point without reprocessing already-consumed messages.

## Image 2: Internals of Message Partitioning

This image demonstrates how Kafka distributes messages across partitions. Here's the explanation:

### 1. Producer Sends Messages:

- The producer sends a batch of 10 messages: `{a, b, c, d, e, f, g, h, i, j}`.

### 2. Partition Assignment:

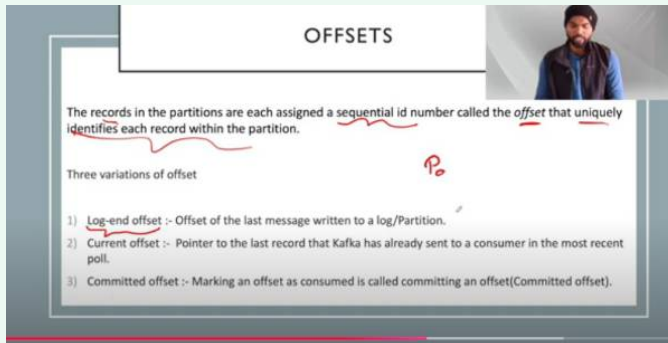
- Messages are assigned to different partitions (P0, P1, P2) in the topic based on the partitioning strategy:
  - Default: Round-robin or based on a key (if provided).
- Messages are distributed as follows:
  - **P0 (Partition 0):** `a, d, g, j`
  - **P1 (Partition 1):** `b, e, h`
  - **P2 (Partition 2):** `c, f, i`

### 3. Offsets in Partitions:

- Each message within a partition is assigned a unique sequential **offset**:
  - **P0:** Offsets `0, 1, 2, 3` correspond to messages `a, d, g, j`.
  - **P1:** Offsets `0, 1, 2` correspond to messages `b, e, h`.
  - **P2:** Offsets `0, 1, 2` correspond to messages `c, f, i`.

## Summary

- The **offset** ensures uniqueness for messages within a partition and helps Kafka track consumed and unconsumed messages.
- The producer distributes messages across partitions to balance the load, following the defined strategy.
- These concepts are critical for ensuring data consistency and efficient message processing in Kafka systems.



Explain this with respect to : Internals of Message Partitioning as attached in picture above



The relationship between **Offsets** (first image) and the **Internals of Message Partitioning** (second image) can be explained as follows:

### Offsets in Kafka:

Offsets are unique sequential IDs assigned to records within a partition. They help Kafka keep track of:

1. The **position of the consumer** (current offset).
2. The **point up to which messages are processed** (committed offset).
3. The **next available position for a new record** (log-end offset).

### Internals of Message Partitioning:

In the second image, messages are distributed across three partitions (**P0, P1, P2**), and each message is assigned an offset within its respective partition.

## Mapping the Concepts:

### 1. Log-End Offset:

- In **P0**, the log-end offset is `3` (the position of message `j`).
- Similarly:
  - For **P1**, the log-end offset is `2` (position of `h`).
  - For **P2**, the log-end offset is `2` (position of `i`).
- The log-end offset marks where new messages will be appended.

### 2. Current Offset:

- This is the offset of the last message the consumer has read.
- If the consumer has read up to `2` in **P0** (message `g`), then the current offset is `2`. The consumer will read `j` next.

### 3. Committed Offset:

- After the consumer processes a message, it commits the offset to ensure that if it crashes, it can resume from the committed position.
- For example:
  - If the consumer commits offset `1` in **P0**, it means messages up to `d` are marked as processed, and processing will restart from offset `2`.

## Flow with Partitioning Example:

- Messages `{a, b, c, d, e, f, g, h, i, j}` are distributed across partitions:
  - **P0**: `a, d, g, j`
  - **P1**: `b, e, h`
  - **P2**: `c, f, i`
- For each partition:
  - The **log-end offset** is where new messages will be appended.
  - The **current offset** tracks the consumer's position during consumption.
  - The **committed offset** ensures fault tolerance by marking consumed records.

## Key Takeaway:

Offsets ensure Kafka tracks the position of each record within a partition. This allows efficient data processing, fault tolerance, and recovery during message consumption while partitioning ensures load distribution across brokers.



Kindly explain above and how does Producer send message. what is the significance of key while sending message



This image demonstrates how a Kafka **Producer** sends messages to a Kafka topic and the significance of a **key** in this process.

## Producer Sends a Message

### 1. Message Structure:

- **Key:** In this case, the key is `null`, meaning no specific partition is assigned based on a key.
- **Payload:** The message payload is `"Brijesh"`, which is the actual data being sent.
- **Topic:** The message is sent to the topic `myTopic`.
- **Partition:** Since the key is `null`, Kafka uses a default partitioning strategy (usually round-robin) to decide the partition. Here, the message is assigned to **Partition P1**.
- **Offset:** Within Partition P1, the message is assigned an offset of `0` (since it is the first message in this partition).

### 2. Flow:

- The producer sends the message to the Kafka cluster.
- Kafka determines the partition for the message (based on the key or round-robin).
- The message is stored in the selected partition with a unique offset.

## Significance of the Key

The **key** plays a critical role in message routing within Kafka. Its significance is as follows:

### 1. Partition Selection:

- If a key is provided, Kafka uses the **hash of the key** to determine the partition.
- Messages with the same key will always go to the same partition, ensuring **order** for messages with the same key.

### 2. Load Balancing:

- If the key is ``null``, Kafka distributes messages evenly across partitions (round-robin strategy). This is useful for **load balancing** when message order isn't important.

### 3. Data Consistency:

- Using a key ensures that related messages (e.g., for a specific user or transaction) are routed to the same partition. This is critical for scenarios where maintaining **message order** is required.

---

## Example in the Image

- Since the key is ``null``, the message ``"Brijesh"`` is routed to **Partition P1** based on Kafka's round-robin logic.
  - It is stored at offset ``0`` in Partition P1.
  - If a key like ``"user123"`` were used, the hash of ``"user123"`` would determine the partition, and all future messages with the same key would go to the same partition, maintaining order for that key.
- 

## Key Takeaways

- The **key** ensures determinism in message routing to partitions, which is useful for maintaining order.
- Without a key, Kafka relies on round-robin partitioning, which distributes the load but doesn't guarantee order across partitions.
- The offset uniquely identifies the message within a partition and is critical for tracking and consuming messages efficiently.



In the message payload what is the significance of timestamp



The **timestamp** in a Kafka message payload holds critical significance in Kafka's messaging system. Here's a detailed explanation of its role:

---

## Significance of the Timestamp

### 1. Message Creation Time:

- The timestamp indicates **when the message was created by the producer**.
- This helps in tracking when the event occurred and can be useful for debugging or auditing purposes.

### 2. Event Ordering:

- Kafka processes messages sequentially within a partition, but when consuming across partitions, timestamps help consumers maintain a logical order of events.

### 3. Time-Based Processing:

- Some applications process messages based on their timestamp rather than the offset. For example:
  - Processing only messages created after a specific time.
  - Running time-windowed aggregations in stream processing.

### 4. Recovery from Failures:

- In case of consumer failure, the timestamp can help resume consumption from messages after a specific point in time, instead of relying on offsets alone.

### 5. Latency Monitoring:

- Comparing the timestamp of a message (when it was produced) with the current time (when it is consumed) helps measure the end-to-end latency in the Kafka pipeline.

### 6. Retention and Expiry:

- Kafka topics can be configured with **time-based retention policies** (e.g., messages older than 7 days are deleted). The timestamp plays a crucial role in determining which messages should be retained or deleted.

---

## Types of Timestamps in Kafka

Kafka provides two types of timestamps:

### 1. Create Time:



- The producer sets this timestamp at the time of message creation.
- It is the default and represents the actual event time.

## 2. Log Append Time:

- Kafka brokers set this timestamp when the message is appended to the log.
- This is useful when the producer's clock is unreliable or unsynchronized.

## Practical Use Cases

### 1. Real-Time Analytics:

- Timestamps allow systems to group, aggregate, or filter messages based on event times.
- For instance, calculating metrics in 1-minute windows requires accurate timestamps.

### 2. Time-Ordered Processing:

- In stream processing (e.g., Kafka Streams, Flink), timestamps are crucial for ordering and windowed computations.

### 3. Replaying Events:

- Consumers can replay events starting from a specific timestamp to debug or reproduce scenarios.

## Example in Context

- If the message `"Brijesh"` has a timestamp `1678901234567`` (Unix epoch format), it signifies:
  - The message was produced at a specific point in time.
  - Consumers can process the message knowing its exact creation time.
  - If offsets are skipped due to failure, consumers can still process messages created after a specific timestamp.

## Key Takeaway

The timestamp enhances Kafka's functionality by enabling time-sensitive operations, debugging, and monitoring, making it essential for many real-world applications, especially in event-driven architectures and streaming use cases.



Is it when the message was written in filesystem regarding timestamp



The **timestamp** in a Kafka message does not necessarily represent the exact time the message was written to the filesystem. Instead, its meaning depends on the **timestamp type** configured for the topic or broker. Here's a detailed breakdown:

## Types of Timestamps in Kafka

### 1. Create Time (Default):

- **Meaning:** The timestamp represents the time the producer created the message.
- **Who Sets It:** The producer sets this timestamp.
- **Use Case:** This is the most common scenario because it reflects the actual event time of the message.
- **Filesystem Aspect:** It does **not** correlate directly to when the message was written to the filesystem or Kafka logs.

### 2. Log Append Time:

- **Meaning:** The timestamp represents the time the Kafka broker **appended the message to the partition's log** (i.e., when it was written to the filesystem).
- **Who Sets It:** The broker sets this timestamp.
- **Use Case:** This is useful when producers' clocks are unreliable, unsynchronized, or when the Kafka system relies on log-append times to enforce order.
- **Filesystem Aspect:** This **does** correlate with the time the message was written to Kafka's underlying storage (filesystem).

## When Is the Timestamp Related to the Filesystem?

- If the **Log Append Time** is configured, the timestamp reflects the time when the broker persisted the message to the partition log on disk.
- This ensures that the timestamp aligns with the message's storage time, which can be important for use cases like log-based monitoring or latency analysis.

## How to Check or Configure Timestamp Type

1. **Broker Level:** Kafka brokers have a default configuration for timestamp types:

- ``log.message.timestamp.type`` can be set to either ``CreateTime`` or ``LogAppendTime``.

2. **Topic Level:** Individual topics can override the broker's default configuration for timestamp type:

- Use ``--config message.timestamp.type=<CreateTime|LogAppendTime>`` when creating or updating a topic.
- 

## Summary

- If the timestamp is set to **Create Time**, it is when the producer created the message, unrelated to the filesystem.
- If the timestamp is set to **Log Append Time**, it represents when the broker wrote the message to the filesystem (Kafka log).

Knowing this distinction is important when designing Kafka pipelines where time-based operations (e.g., retention, ordering, or time-window processing) are critical.